

CRE rental-market intelligence — a demand-side early-warning engine, validated across 17 markets.

79.3%

DETECTION HIT RATE
WALK-FORWARD, NO
LOOKAHEAD

17

MARKETS VALIDATED
4 TIERS · 58
CONFIRMED TURNS

5–7mo

DEMAND → RENT LEAD
MEDIAN 5MO
DOMINANT

16/17

POSTINGS CLEAN
JOLTS QUILTS 14/17

01 — THE PROBLEM

Mid-market debt managers react to rent — they don't lead it.

CRE rent forecasting runs on proprietary, paywalled data — CoStar, RealPage, Yardi — and the space is consolidating fast. Mid-market debt managers underwriting floating-rate portfolios have no systematic way to monitor the demand signals that actually move rent. They react to rent changes after they show up in the index, by which point the decision is already made.

02 — THE ENGINE

Demand leads rent by 5–7 months. The engine detects the turn.

Job postings, JOLTS quit-rates, WARN layoff notices, and employment turn 5–7 months before rent growth does. Walk-forward backtest: **79.3% hit rate across 17 markets** (46 of 58 confirmed turns). The result held byte-identical when we doubled the sample from 9 to 17 metros — a fitted model degrades when you add data; this didn't move. Postings is the universal leader; JOLTS quits clean in 14/17; WARN is the multi-cycle depth signal.

03 — THE PRODUCT

Three layers, the way an analyst would actually use it.

01

Full market analysis

All 17 metros at a glance — each market's confirmed directional call, dominant signal, and live concession read.

02

Agentic drill-down

Click any market and ask the analyst agent why it's firming or softening. Every quantitative answer comes from a validated tool — never the model's invention — with sources on every claim.

03

Backtesting validation

Signal strength is measured, not asserted — hit rate, skill versus baseline, and an honest rotation backtest, all queryable.

04 — WHAT WE REMOVED

The discipline is the product. We tested every public leading indicator against pre-stated criteria — and dropped what failed.

Permits**DROPPED · PRO-CYCLICAL**

Co-moves with rent rather than leading it. Builders permit *when* demand fires.

Wages**DROPPED · WRONG-SIGNED 3 OF 9**

Adding it dropped the hit rate from 79.3% to 71.9%. BSS vs base rate collapsed from +0.31 to +0.05.

Rent-vs-own ratio**DROPPED · MECHANICAL CORRELATION**

Failed in 8 of 9 markets. ZORI sits in both the denominator and the target — the ratio responds to rent rather than leads it.

We also tested against a persistence baseline most models skip — and learned the engine's real job. Trading the signal underperforms: in illiquid CRE you can be right on the fundamentals and still get caught by the rate cycle. This is an **early-warning tool for avoid and build decisions, not a trading signal**. Saying so is the credibility.

05 — BRIGHT DATA: THE LIVE LAYER

Two anti-bot sources scraped live. One convergent answer.

Two demand-texture layers no public API exposes — scraped through **Bright Data Web Unlocker** from LinkedIn job postings and apartments.com concessions. A live public-data proxy for signals institutions license from CoStar / RealPage.

SUN BELT — HIGH CONCESSION SHARE

engine flagged these on labor fundamentals

Atlanta	90%	
Phoenix	88%	
Austin	83%	

COASTAL CONSTRAINED — LOW CONCESSION

supply-tight, landlord-favored

San Francisco	28%	
New York	24%	
Boise	15%	

Two independent data sources, two different methods, the same answer. *The live layer corroborates the engine.*

LIVE TODAY

- 17-market engine
- Agentic drill-down with sources
- Live Bright Data concessions + postings
- Backtest validation

ROADMAP

- Concession history once retained
- Transaction-price validation (RCA / Green Street)
- Longer-hold backtests
- Licensed lease-trade-out signal

APP

homestardata.com

GITHUB

github.com/Masmedeam/groundswell

DEMO VIDEO

[TBD]